

EE379K: Lab 2

Group Members: Irfan Hasan (ih3976), Peter Zhang (yz7724)

Question 1

```
In [34]: import pandas as pd
import seaborn as sns
import numpy as np
from pandas.plotting import scatter_matrix
import matplotlib.pyplot as plt

df = pd.read_csv('Lab2_Data/DF1')
df = df.drop(df.columns[[0]], axis=1)
corr = df.corr()
print("-----1a-----")
print("-----Correlation coefficients from pandas-----\n")
print(corr)

print("\n----- Heatmap from Seaborn ----- \n")
sns.heatmap(corr)
plt.show()

print("\nFrom the data it can be seen that the following columns are correlated")
print("(0,2), (1,3)\n")

print("----- 1b -----")
print("Covariance matrix is the pairwise covariance between all the columns")
print(df.cov())
print("The covariance matrix reflects the results seen in the plots. Specifically")

print("\n----- 1c -----")
cov = [[3, 0, 0], [0, 1.5, 0.5], [0, 0.5, 6]]
print('Chosen covariance:\n{}\n'.format(np.matrix(cov)))
samples = [50, 100, 500, 1000, 2500, 5000, 10000, 100000]
res = []

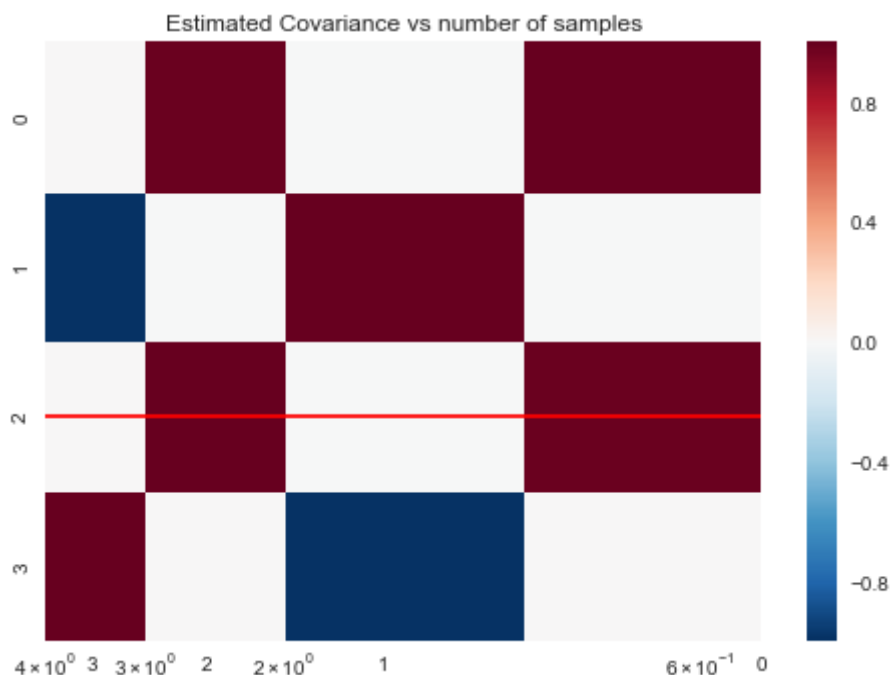
for n in samples:
    sample = np.random.multivariate_normal([0,0,0], cov, n)
    estimated_cov = np.cov(sample, rowvar=False)
    res.append(estimated_cov[1][1])

fig = plt.figure()
ax = plt.gca()
ax.scatter(samples, res)
ax.set_xscale('log')
ax.set_title('Estimated Covariance vs number of samples')
plt.axhline(y=1.5, c='r')
plt.xlabel('Sample')
plt.ylabel('Estimated Covariance')
plt.show()
```

```
-----1a-----
-----Correlation coefficients from pandas-----
```

	0	1	2	3
0	1.000000	-0.003998	0.990066	0.004111
1	-0.003998	1.000000	-0.004085	-0.990235
2	0.990066	-0.004085	1.000000	0.004067
3	0.004111	-0.990235	0.004067	1.000000

----- Heatmap from Seaborn -----



From the data it can be seen that the following columns are correlated:
(0,2), (1,3)

----- 1b -----

Covariance matrix is the pairwise covariance between all the columns in the dataset

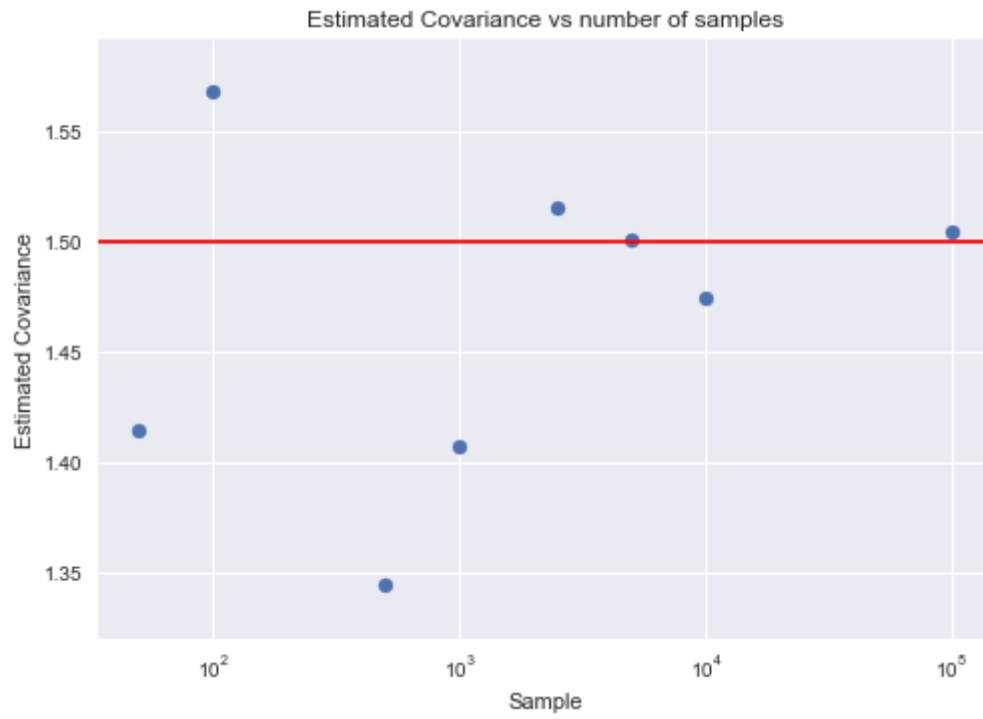
	0	1	2	3
0	1.001558	-0.004012	0.991624	0.004125
1	-0.004012	1.005378	-0.004099	-0.995457
2	0.991624	-0.004099	1.001589	0.004081
3	0.004125	-0.995457	0.004081	1.005168

The covariance matrix reflects the results seen in the plots. Specifically we can see that columns 0 and 2 and 1 and 3 have the highest covariance.

----- 1c -----

Chosen covariance:

```
[[ 3.  0.  0. ]
 [ 0.  1.5 0.5]
 [ 0.  0.5 6. ]]
```



Question 2

```
In [22]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler
%matplotlib inline

print('----- Original plot -----\\n')
df2 = pd.read_csv('Lab2_Data/DF2')
df2 = df2.ix[:, 1:]
df2.plot.scatter(x='0', y='1')
plt.show()

scaler = MinMaxScaler()
df2_scale = scaler.fit_transform(df2)
df2_scaled = pd.DataFrame(data=df2_scale)


print('\\nWe used the MinMaxScaler as it would shrink the x and y axes to a
print('We believe shrinking the x-axis especially will reveal the outlier wh
print('\\n----- Transformed plot -----\\n')
g = df2_scaled.plot.scatter(x=0, y=1)
g.set_ylim([0, 1])
plt.show()
```

----- Original plot -----

```
/Users/irfanhasan/anaconda/lib/python3.6/site-packages/ipykernel_launcher
r.py:7: DeprecationWarning:
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:


```
http://pandas.pydata.org/pandas-docs/stable/indexing.html#deprecate\_ix (h
http://pandas.pydata.org/pandas-docs/stable/indexing.html#deprecate\_ix)
import sys
```



We used the MinMaxScaler as it would shrink the x and y axes to a range from 0 to 1.

We believe shrinking the x-axis especially will reveal the outlier which is more prominent on the y-axis.

----- Transformed plot -----



Question 3

```
In [1]: import numpy as np

def calc_std_dev(n):
    deltas = []
    for i in range(n):
        X = np.random.randn(n)
        E = np.random.randn(n)
        y = -3 + np.dot(X, 0) + E
        beta_h = np.dot(X,y) / np.dot(X,X)
        deltas.append(beta_h)
    return np.std(deltas)

calc_std_dev(150)
```

Out[1]: 0.24904452509546193

We can see that $B_{\text{hat}} = -0.15$ is not as significant since the empirical standard deviation of the error is much larger than 0.15, so the error accounts for all of it.

```

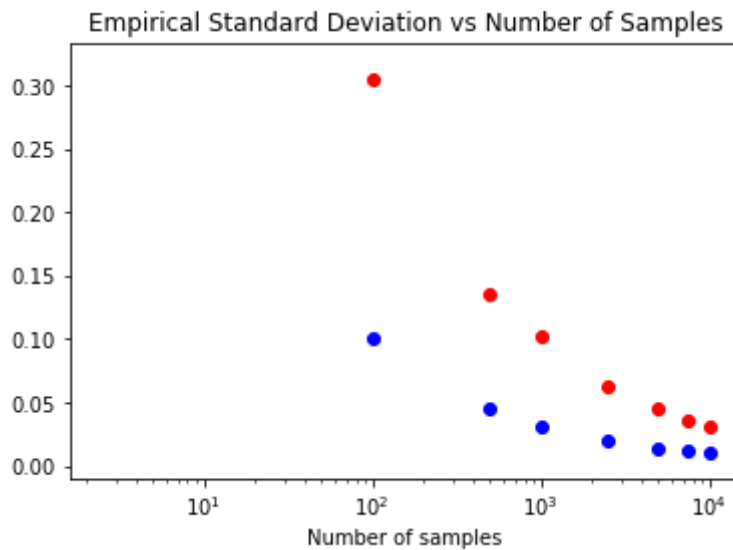
In [5]: import math
import numpy as np
import matplotlib.pyplot as plt

def calc_std_dev(n):
    deltas = []
    for i in range(n):
        X = np.random.randn(n)
        E = np.random.randn(n)
        y = -3 + np.dot(X, 0) + E
        beta_h = np.dot(X,y) / np.dot(X,X)
        deltas.append(beta_h)
    return np.std(deltas)

samples = [100, 500, 1000, 2500, 5000, 7500, 10000]
std_devs = []
one_over = []
for n in samples:
    std_dev = calc_std_dev(n)
    std_devs.append(std_dev)
    one_over.append(1/math.sqrt(n))

fig = plt.figure()
ax = plt.gca()
ax.scatter(samples, std_devs, c='r')
ax.scatter(samples, one_over, c='b')
ax.set_xscale('log')
ax.set_title('Empirical Standard Deviation vs Number of Samples')
plt.xlabel('Number of samples')
plt.show()
print('The fit is good.')

```



The fit is good.

Question 4

```

In [33]: import sys
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

class Question4:
    def get_k_names(self, k, year):
        filename = "Names/yob" + str(year) + ".txt"
        data = pd.read_csv(filename, sep=",", header=None)
        print data.head(k)

    def name_frequency(self, name):
        m = 0
        f = 0
        for year in range(1880, 2016):
            filename = "Names/yob" + str(year) + ".txt"
            data = pd.read_csv(filename, sep=",", header=None)
            data = data[data[0] == name]
            for row in data.itertuples():
                if row[2] == 'M':
                    m += row[3]
                else:
                    f += row[3]

        print "For name " + name
        print "Male: " + str(m)
        print "Female: " + str(f)

    def relative_frequency(self, name, year):
        filename = "Names/yob" + str(year) + ".txt"
        data = pd.read_csv(filename, sep=",", header=None)
        total = data[2].sum()
        data = data[data[0] == name]

        print "For year " + str(year)
        for row in data.itertuples():
            print "{0} {1} {2:.9f}".format(row[1], row[2], float(row[3])/total)

    def change_in_pop(self):
        result = set()
        names = dict(dict()) # {name : []}

        for year in range(1880, 2016):
            filename = "Names/yob" + str(year) + ".txt"
            data = pd.read_csv(filename, sep=",", header=None)

            for row in data.itertuples():
                if row[1] not in names:
                    names[row[1]] = {}
                if year not in names[row[1]]:
                    names[row[1]][year] = 0

                if row[2] == 'M':
                    names[row[1]][year] += row[3]
                else:
                    names[row[1]][year] -= row[3]

```



```

for name, entries in names.iteritems():
    pos = neg = False
    for y in sorted(entries.iterkeys()):
        if entries[y] > 0 and neg:
            result.add(name)
            break
        elif entries[y] < 0 and pos:
            result.add(name)
            break
        elif entries[y] > 0:
            pos = True
        elif entries[y]:
            neg = True

for n in result:
    print n

```

```
q4 = Question4()
```

Write a program that on input k and XXXX, returns the top k names from year XXXX

```
In [5]: q4.get_k_names(5, 1996)
```

```

      0  1      2
0      Emily  F  25150
1      Jessica F  24192
2      Ashley F  23676
3      Sarah  F  21029
4      Samantha F  20545

```

Write a program that on input Name returns the frequency for men and women of the name Name

```
In [9]: q4.name_frequency('Bailey')
```

```

For name Bailey
Male: 20457
Female: 91648

```

Modify the above program to return relative frequency.

```
In [15]: q4.relative_frequency('Bailey', 1996)
```

```

For year 1996
Bailey F 0.001149507
Bailey M 0.000425134

```

Find all names that used to be more popular for one gender, but then became more popular for another gender.

```
In [ ]: q4.change_in_pop() # names are not printed due to there are too many of them
```

Question 5

Tutorial by Dataquest

```
In [2]: import pandas as pd

tweets = pd.read_csv("tweets.csv")
tweets.head()
```

```
Out[2]:
```

		id	id_str	user_location	user_bg_color	retweet_count	user_name	polarity	
0	1	729828033092149248		Wheeling WV	022330	0	Jaybo26003	0.00	101
1	2	729828033092161537		NaN	C0DEED	0	britttany_ns	0.15	101
2	3	729828033566224384		NaN	C0DEED	0	JeffriesLori	0.00	101
3	4	729828033893302272		global	C0DEED	0	WhorunsGOVs	0.00	101
4	5	729828034178482177		California, USA	131516	0	BJCG0830	0.00	101


```
In [3]: def get_candidate(row):
    candidates = []
    text = row["text"].lower()
    if "clinton" in text or "hillary" in text:
        candidates.append("clinton")
    if "trump" in text or "donald" in text:
        candidates.append("trump")
    if "sanders" in text or "bernie" in text:
        candidates.append("sanders")
    return ",".join(candidates)

tweets["candidate"] = tweets.apply(get_candidate,axis=1)
```

```
In [5]: import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
```

```
In [6]: counts = tweets["candidate"].value_counts()
plt.bar(range(len(counts)), counts)
plt.show()

print(counts)
```




```
trump          119998
clinton,trump   30521
               25429
sanderson      25351
clinton        22746
clinton,sanders 6044
clinton,trump,sanders 4219
trump,sanders  3172
Name: candidate, dtype: int64
```

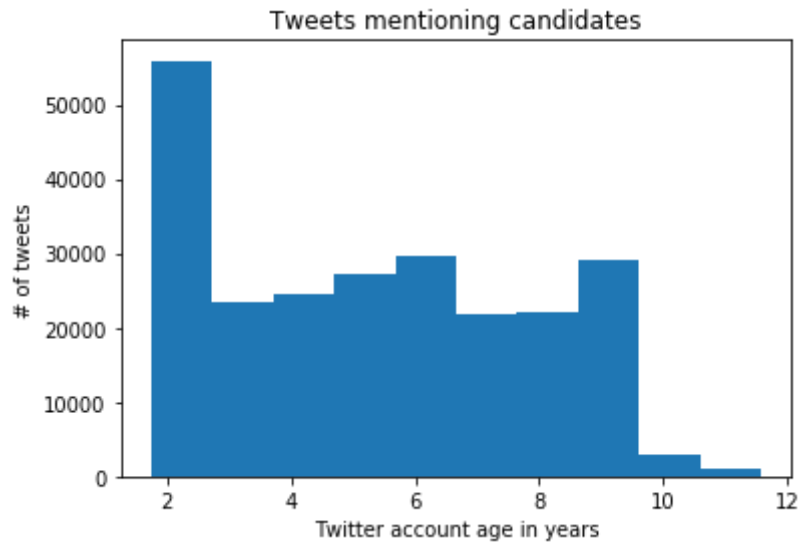
```
In [7]: from datetime import datetime

tweets["created"] = pd.to_datetime(tweets["created"])
tweets["user_created"] = pd.to_datetime(tweets["user_created"])


tweets["user_age"] = tweets["user_created"].apply(lambda x: (datetime.now()
plt.hist(tweets["user_age"])
plt.show()
```



```
In [8]: plt.hist(tweets["user_age"])
plt.title("Tweets mentioning candidates")
plt.xlabel("Twitter account age in years")
plt.ylabel("# of tweets")
plt.show()
```



```
In [14]: cl_tweets = tweets["user_age"][tweets["candidate"] == "clinton"]
sa_tweets = tweets["user_age"][tweets["candidate"] == "sanderson"]
tr_tweets = tweets["user_age"][tweets["candidate"] == "trump"]
plt.hist([
    cl_tweets,
    sa_tweets,
    tr_tweets
],
    stacked=True,
    label=["clinton", "sanderson", "trump"]
)
plt.legend()
plt.title("Tweets mentioning each candidate")
plt.xlabel("Twitter account age in years")
plt.ylabel("# of tweets")
plt.show()
```



```
In [30]: import matplotlib.colors as colors

tweets["red"] = tweets["user_bg_color"].apply(lambda x: colors.hex2color('#f08080'))
tweets["blue"] = tweets["user_bg_color"].apply(lambda x: colors.hex2color('#1e90ff'))
```

```
In [31]: fig, axes = plt.subplots(nrows=2, ncols=2)
ax0, ax1, ax2, ax3 = axes.flat


ax0.hist(tweets["red"])
ax0.set_title('Red in backgrounds')

ax1.hist(tweets["red"][tweets["candidate"] == "trump"].values)
ax1.set_title('Red in Trump tweeters')

ax2.hist(tweets["blue"])
ax2.set_title('Blue in backgrounds')

ax3.hist(tweets["blue"][tweets["candidate"] == "trump"].values)
ax3.set_title('Blue in Trump tweeters')

plt.tight_layout()
plt.show()
```



```
In [32]: tweets["user_bg_color"].value_counts()
```

```
Out[32]: C0DEED      108977
          000000      31119
          F5F8FA      25597
          131516       7731
          1A1B1F       5059
          022330       4300
          0099B9       3958
          642D8B       3767
          FFFFFFFF      3101
          9AE4E8       2651
          ACDED6       2383
          352726       2338
          C6E2EE       1978
          709397       1518
          EBEBEB       1475
          FF6699       1370
          BADFCD       1336
          FFF04D       1300
          EDECE9       1225
          B2DFDA       1218
          DBE9ED       1113
          ABB8C2       1101
          8B542B       1073
          3B94D9        623
          89C9FA        414
          DD2E44        351
          94D487        318
          4A913C        300
          9266CC        287
          F5ABB5        267

          ...
          5470A8         1
          00AEFF         1
          C49C4B         1
          778877         1
          09380E         1
          09536E         1
          3D3C3D         1
          48394D         1
          3D3C3A         1
          140C0E         1
          AE1BCF         1
          EBE39B         1
          056785         1
          FCF3EA         1
          2E332F         1
          FCF7F8         1
          FCF7F7         1
          0F6B2C         1
          1D1F1B         1
          180018         1
          2686B3         1
          8F0E8F         1
          CCD4E8         1
          FFEF42         1
```

```

08F5F5      1
4E5254      1
42373E      1
272D29      1
F00CC2      1
A3004D      1
Name: user_bg_color, Length: 6970, dtype: int64

```

```
In [33]: tc = tweets[~tweets["user_bg_color"].isin(["CODEED", "000000", "F5F8FA"])]
```

```

def create_plot(data):
    fig, axes = plt.subplots(nrows=2, ncols=2)
    ax0, ax1, ax2, ax3 = axes.flat

    ax0.hist(data["red"])
    ax0.set_title('Red in backgrounds')

    ax1.hist(data["red"][data["candidate"] == "trump"].values)
    ax1.set_title('Red in Trump tweets')


    ax2.hist(data["blue"])
    ax2.set_title('Blue in backgrounds')

    ax3.hist(data["blue"][data["candidate"] == "trump"].values)
    ax3.set_title('Blue in Trump tweeters')

    plt.tight_layout()
    plt.show()

create_plot(tc)

```





```
In [34]: gr = tweets.groupby("candidate").agg([np.mean, np.std])

fig, axes = plt.subplots(nrows=2, ncols=1, figsize=(7, 7))
ax0, ax1 = axes.flat

std = gr["polarity"]["std"].iloc[1:]
mean = gr["polarity"]["mean"].iloc[1:]
ax0.bar(range(len(std)), std)
ax0.set_xticklabels(std.index, rotation=45)
ax0.set_title('Standard deviation of tweet sentiment')

ax1.bar(range(len(mean)), mean)
ax1.set_xticklabels(mean.index, rotation=45)
ax1.set_title('Mean tweet sentiment')

plt.tight_layout()
plt.show()
```



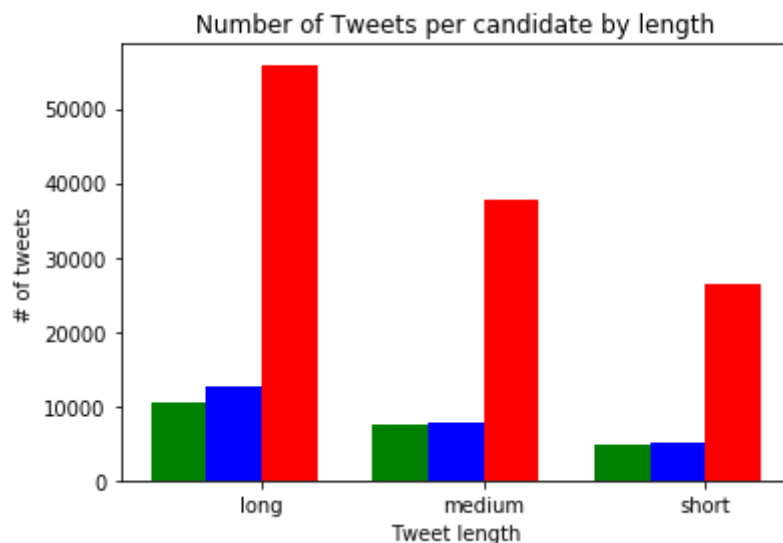
```
In [38]: def tweet_lengths(text):
    if len(text) < 100:
        return "short"
    elif 100 <= len(text) <= 135:
        return "medium"
    else:
        return "long"

    tweets["tweet_length"] = tweets["text"].apply(tweet_lengths)

    tl = {}
    for candidate in ["clinton", "sanderson", "trump"]:
        tl[candidate] = tweets["tweet_length"][tweets["candidate"] == candidate]
```

```
In [39]: fig, ax = plt.subplots()
    width = .5
    x = np.array(range(0, 6, 2))
    ax.bar(x, tl["clinton"], width, color='g')
    ax.bar(x + width, tl["sanderson"], width, color='b')
    ax.bar(x + (width * 2), tl["trump"], width, color='r')

    ax.set_ylabel('# of tweets')
    ax.set_title('Number of Tweets per candidate by length')
    ax.set_xticks(x + (width * 1.5))
    ax.set_xticklabels(('long', 'medium', 'short'))
    ax.set_xlabel('Tweet length')
    plt.show()
```



Aggregate the results by state.

```
In [14]: '''
Adding to the filters for each state will increaes the number of captures
'''

filters = [
    ['al', 'alabama'],
    ['ak', 'alska'],
    ['az', 'arizona'],
    ['ar', 'arkansas'],
    ['ca', 'cali', 'california'],
    ['co', 'colorado'],
    ['ny'],
    ['pa', 'pittsburgh'],
    ['tx', 'texas', 'austin', 'houstin'],
    ['va', 'virginia'],
    ['wv'],
    ['wy']]

def get_state(row):
    result = []
    location = str(row).lower().split(' ')

    found = False
    for word in location:
        found = False
        for f in filters:
            for addr in f:
                if addr == word:
                    found = True
                    result.append(f[0])
                    break
            if found:
                break
        if found:
            break

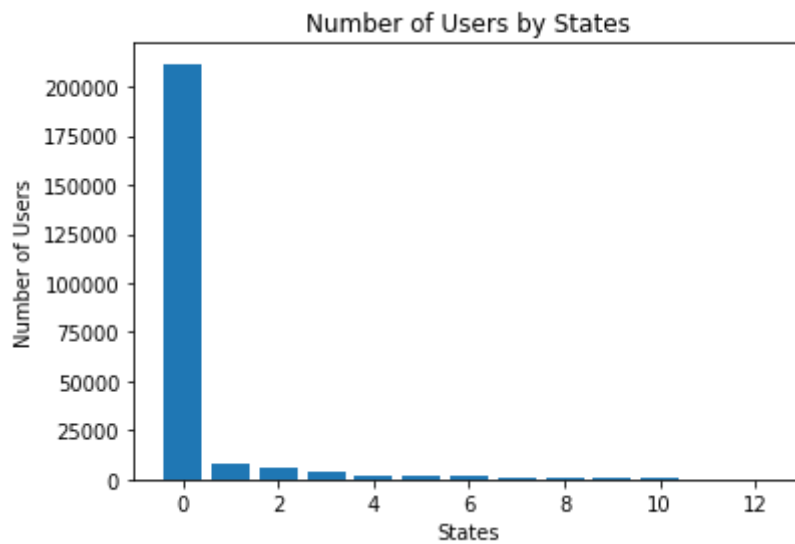
    if found == False:
        result.append('N/A')

    return ", ".join(result)

tweets["state"] = tweets['user_location'].apply(get_state)
```

```
In [15]: counts = tweets['state'].value_counts()
plt.bar(range(len(counts)), counts)
plt.title("Number of Users by States")
plt.xlabel("States")
plt.ylabel("Number of Users")
plt.show()

print (counts)
```



```
N/A      211669
ca        8394
...
```

Written Questions

1.

$$a. \quad Z_{avg} = \sum_{i=1}^n \frac{Z_i}{n}$$

when $n=10,000$

$$\mu_{avg} = 0 \quad \sigma_{avg} = \frac{1}{\sqrt{10,000}} = 0.01$$

$$\text{Thus } Z_{avg} \sim N(0, 0.01^2)$$

Using normal table

$$P(Z_{avg} > 0.1) = 0$$

$$P(Z_{avg} > 0.01) = 0.1587$$

$$P(Z_{avg} > 0.001) \approx 0.4602$$

$$b. \quad Z \sim N(\mu, \sigma^2)$$

$$\mu_{avg} = \mu$$

$$\sigma_{avg} = \frac{\sigma}{\sqrt{n}}$$

$$Z_{avg} \sim N(\mu, \frac{\sigma^2}{n})$$

 $Z_{avg} - \mu$ shifts Z_{avg} 's

mean to 0.

$$i. \quad \text{Thus } Z = \frac{x - \mu}{\sigma}$$

$$P(Z_{avg} - \mu > n^{-1/3}) = \int_{\frac{-1/3}{\sigma}}^{\infty} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx$$

$$ii. \quad P(Z_{avg} - \mu > n^{-1/2}) = \int_{\frac{-1/2}{\sigma}}^{\infty} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx$$

$$iii. \quad P(Z_{avg} - \mu > n^{-2/3}) = \int_{\frac{-2/3}{\sigma}}^{\infty} \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx$$

EE379K HW #2

Question #2

$$X_i^2 \beta^2 - 2X_i y_i \beta + y_i^2$$

$$a) \frac{1}{n} \sum_{i=1}^n (X_i \beta - y_i)^2$$

$$\frac{1}{n} [(X_1 \beta - y_1)^2 + (X_2 \beta - y_2)^2 + \dots + (X_n \beta - y_n)^2]$$

$$\frac{1}{n} [\beta^2 (X_1^2 + X_2^2 + \dots + X_n^2) - 2\beta (X_1 y_1 + X_2 y_2 + \dots + X_n y_n) + (y_1^2 + \dots + y_n^2)]$$

$$\frac{1}{n} [\beta^2 \sum_{i=1}^n X_i^2 - 2\beta \sum_{i=1}^n X_i y_i + \sum_{i=1}^n y_i^2]$$

$$A = \frac{1}{n} \sum_{i=1}^n X_i^2 \quad B = -\frac{2}{n} \sum_{i=1}^n X_i y_i \quad C = \sum_{i=1}^n y_i^2$$

$A \geq 0$ because the value is squared and any number squared is non-negative.

$$b) \frac{d}{d\beta} \min_{\beta} = \frac{2}{n} \beta \sum_{i=1}^n X_i^2 - \frac{2}{n} \sum_{i=1}^n X_i y_i$$

$$\beta \sum_{i=1}^n X_i^2 = \sum_{i=1}^n X_i y_i$$

$$\hat{\beta} = \frac{\sum_{i=1}^n X_i y_i}{\sum_{i=1}^n X_i^2}$$

$$c) \hat{\beta} = \frac{\sum X_i (X_i \beta + e_i)}{\sum X_i^2} \quad \hat{\beta} = \beta + \frac{\sum X_i e_i}{\sum X_i^2}$$

$$\hat{\beta} = \frac{\sum X_i^2 \beta + \sum X_i e_i}{\sum X_i^2} \quad Z e = \frac{\sum X_i e_i}{\sum X_i^2}$$

$$\hat{\beta} = \frac{\sum X_i^2 \beta}{\sum X_i^2} + \frac{\sum X_i e_i}{\sum X_i^2} \quad Z = \frac{\vec{X}}{\|\vec{X}\|^2}$$